

An Autonomous ML Research Agent for KuaiRand-Pure

System Architecture

TikTok TechJam 2026, Track 2

Abstract

This document describes the architecture of an LLM-driven agent that runs an end-to-end machine learning research loop without human intervention: it inspects a dataset it has never seen, reproduces a published baseline, then repeatedly forms a hypothesis, writes the training script, executes it in a sandbox, verifies the result against a trusted evaluator, revises what it believes, and chooses what to try next. The design question it answers is not “can a model write a training script” but “what has to surround the model so that an unattended run produces a result you can trust.” Most of the system is that surrounding machinery: a trust boundary between the controller and model-written code, an evaluator that never believes a script’s own reported score, an integrity critic, a search tree with retirement, and a revisable belief set. Section 9 reports the submitted run, **r96**, in full: it reached the organizers’ convergence rule in 14 of the 50 permitted iterations, in 53 minutes on CPU alone, with zero failures and zero human interventions, for a validation PRIMARY of 0.605759 against the official baseline’s 0.6016. Section 10 draws out what the run establishes — chiefly that charging the budget per *script* rather than per model let it compare 316 candidates for 14 metered iterations. Section 11 gives the constraint-by-constraint position against the problem statement, and Section 12 records how each reported figure was checked.

1 Overview

The task is within-user ranking of logged short-video impressions from the KuaiRand-Pure dataset. The scored label is the native column `long_view`; the metric is

$$\text{PRIMARY} = \frac{1}{2}(\text{GAUC} + \text{nDCG@5}),$$

with GAUC the positive-count-weighted mean of per-user AUC over users satisfying $0 < \text{positives} < \text{impressions}$, and nDCG@5 averaged over *all* users, so a user with no positive contributes zero. Splits are fixed by date: train 20220408–20220421 (1,141,112 rows), validation 20220422–20220428 (124,909), hidden test 20220429–20220508 (170,588).

The agent is dataset-agnostic in construction: everything it knows about KuaiRand is either measured from the cache at start-up (`agent/facts.py`) or established by its own experiments. The system is roughly 12,100 lines of Python with no third-party dependency in the controller itself — the LLM client is written against `urllib` — plus a 4,300-line React viewer that reads the run records off disk.

The training-data rule. One constraint runs through the whole design and is enforced at four separate layers, so it is worth stating once up front. Models may be fit on the **train split only** (20220408–20220421), for *both* the validation and the test predictions they save. Validation is feedback, not a second training set, and that extends beyond weights to anything *chosen* by watching validation — an early-stopping round count, a threshold, a blend weight. The four layers are: the proposer’s RULES block states it; the critic detects violations by AST (Section 5); `data_contract` in `run_meta.json` records which contract a run ran under; and the run viewer classifies runs under a superseded contract as **legacy-contract** rather than counting them.

1.1 Design constraints that shaped it

Four scored criteria drive most of the non-obvious decisions:

- **Autonomy** is measured by intervention count, so every failure mode has to be handled in-loop. A crashed script’s traceback is returned to the proposer with its own source; a timeout is fed back differently, because a timeout is not a bug to fix. Ideas are retired rather than escalated.
- **Cost** is scored, so the LLM is kept out of the training loop entirely. It is called twice per iteration in the sequential configuration — once to propose, once to revise beliefs — and everything else is deterministic Python. Prompt history is summarised to a fixed window so per-iteration token use stays flat as a run grows.
- **Feasibility** is scored on wall-clock, not accumulated script time, so the loop measures and enforces true elapsed time against a six-hour backstop.
- **Innovation** is scored on what *the agent* identified as worth trying, which is why the task brief is restricted to specification and API — what to try has to come from the agent’s own measurements. Section 10 traces that chain on the submitted run, from the drift its EDA measured to the methods it chose in response.

2 Repository layout

Module	Role	LOC	Trust
<i>agent/ — the product</i>			
loop.py	controller; the three phases, budgets, turn loop	955	trusted
proposer.py	prompt construction and response parsing	623	trusted
llm.py	stdlib LLM client, recording, replay	393	trusted
portfolio.py	slots, archive, refill, correlation gate	288	trusted
memory.py	distils prior run ledgers into context	229	trusted
knowledge.py	the revisable belief set	225	trusted
consultant.py	one cross-slot synthesis per turn	203	trusted
ensemble.py	incumbent retention and rank blending	196	trusted
tree.py	search over solution scripts	173	trusted
recovery.py	retry, then retire the idea	136	trusted
critic.py	integrity gate on a scored result	154	trusted
executor.py	sandboxed subprocess, credential stripping	110	boundary
selection.py	user-grouped fold split, accept/reject	108	trusted
diagnose.py	per-segment error profile, drift report	99	trusted
facts.py	measures dataset facts from the cache	86	trusted
kb.py	keyword retrieval over the paper catalogue	57	trusted
ledger.py	append-only JSONL, one record per iteration	52	trusted
<i>pipeline/ — the sandbox the agent works in</i>			
data.py	loader, fixed date splits, train-only vocabs	525	trusted
evaluate.py	GAUC / nDCG@5, pinned to the official metric	202	trusted
history.py	leakage-safe train-only entity histories	109	trusted
submit.py	submission writer and validator	141	trusted

Table 1: Principal modules. *Boundary* marks the single place where trusted control passes into model-written code. `report_run.py` (488) and the `viz/` React app read finished runs; `research/` holds human analysis explicitly off the submission path.

Reading a finished run. `viz/` is a local React application over the run records. Its landing view is derived entirely from `runs/`, never from the prose write-ups, so it cannot drift away from what the runs recorded. Its most useful piece is `run-index.ts`, which *classifies* every run folder from its own metadata rather than from a hand-maintained list: `submitted` (the run’s `submission.csv` is sha256-identical to `submission_best.csv`), `eligible`, `leakage` (`api_surface` exposed full-month item outcome columns), `legacy-contract`, `bonus-dataset`, `unverified` (no `run_meta.json`, so the contract cannot be checked), and `scratch`. `unverified` is deliberately not a verdict — those runs are neither claimed nor excluded, because the records alone cannot settle it.

3 The control loop

`run_agent.py` builds the collaborators — proposer, evaluator, ledger, tree, recovery, knowledge — and hands them to `agent.loop.run_loop`, which owns the entire run. Figure 1 shows one iteration end to end.




Figure 1: One iteration. The arrow leaving the sandbox carries *arrays, never a verdict*: a generated script is part of the search space, not part of the evaluator. Nothing reaches the ledger, the tree or the submission until the score has been recomputed, blended by the controller and passed by the critic.

3.1 Phases

`run_loop` walks three phases in order, choosing on each pass of the `while` loop (`loop.py:534–545`):

1. **eda** — the agent writes and runs its own inspection script. Judged only on whether it ran and printed something; there is no score. Its stdout is persisted to `eda_report.txt` and injected into every later prompt. This is the agent’s *only* source of dataset knowledge. Two attempts.
2. **baseline** — the agent stands up an end-to-end Factorization Machine pipeline. Accepted only if it lands within 0.002 of the organizers’ published validation PRIMARY = 0.6016. A score far *above* the baseline is fed back as a suspected leak, a score far below as a pipeline error. On success the script becomes the root of the search tree. Four attempts, after which the run stops with `baseline_not_reproduced`.
3. **improve** — the turn loop of Figure 1, until a stop condition fires.

A run terminates on the first of: convergence, the 50-script cap, the six-hour wall-clock backstop, an exhausted proposer, a daily LLM quota, or a broken environment (five consecutive sub-second failures with no output at all).

3.2 The convergence rule

The organizers specify that a run has converged when the validation score “has not improved by more than $\varepsilon = 0.002$ over the last $N = 3$ consecutive iterations.” This is read as written — the gain *across* the window, not each single step clearing ε (`loop.py:479`):

$$\text{converged} \iff |c| > N \wedge c_{-1} - c_{-1-N} \leq \varepsilon,$$

where c is the running best-validation curve. The distinction is material: a run climbing +0.0008 per iteration has improved by 0.0024 over three and is not converged, whereas the per-step reading would stop it. The stricter reading is nonetheless recorded as `strict_convergence_iteration` in `run_meta.json`, so a run can be checked against either. Both ε and N are also written to the metadata, so a diagnostic run launched with `--patience 999` is distinguishable from a default run that simply never converged.

4 Prompt construction

`proposer.py` assembles the prompt from ordered blocks. The ordering is not cosmetic: providers discount input that repeats as a stable *prefix*, so blocks that are identical across a run — the task brief, the paper catalogue, the parent script — are placed first, and blocks rewritten every iteration follow. The same text is sent either way; only the cache hit rate differs.

The response contract is one **HYPOTHESIS**: line and exactly one **PYTHON** fence. A reply that fails to parse, or whose hypothesis matches a retired idea, is retried once with a corrective note; a second failure returns `None` and the run stops with `exhausted`.

5 The trust boundary

The generated script is treated as adversarial in three distinct ways, because the three risks are independent.

Credentials. Model-written code runs unsandboxed with the controller’s environment. A measured run exposed `OPENAI_API_KEYS`, `GITHUB_PACKAGE_TOKEN` and others to the child. `executor._strip_credentials` therefore denies by pattern (`KEY`, `TOKEN`, `SECRET`, `AUTH`, ...) rather than allowing by name: an allowlist of what a training script needs is impossible to keep

Block	Contents
TASK_BRIEF	Metric definitions, split row counts and date windows, the pipeline API signature by signature, the output contract, the training-data rule, the environment. Numbers are interpolated from <code>agent/facts.py</code> , measured from the cache, so pointing the harness at another KuaiRand variant does not feed it false premises.
phase instruction	<code>_EDA_TASK</code> , <code>_BASELINE_TASK</code> or <code>_IMPROVE_TASK</code> .
literature catalogue	All 33 entries, one line each — “it lists every method available, including ones that may not suit this data; deciding which are relevant is your job.”
EDA output	The agent’s own inspection results, truncated to 4000 characters.
parent script	The chosen node’s <i>actual source</i> , in full (14 000 characters). This is what makes iteration $N+1$ able to build on N rather than restarting from a skeleton.
diagnosis	<code>diagnose.segment_report</code> : the incumbent’s nDCG@5 against the per-user ceiling, bucketed by impression count, with each bucket’s share of the total gap.
history, beliefs, blacklist	The last 8 iterations relabelled by what happened to the <i>score</i> rather than the loop status; the current belief set; retired ideas.
mode instruction	<code>_SWEEP_</code> , <code>_TUNE_</code> , <code>_BROADEN_INSTRUCTION</code> .
siblings, seed note	Portfolio only: what other slots are attempting this turn; why a revived slot is resuming.

Table 2: Prompt blocks, in assembly order.

correct across platforms, and getting the denylist slightly wide only costs a variable nothing needed.

The hidden test set. `AGENT_HIDE_TEST_LABELS=1` makes `load("test").y` return a `HiddenTestLabels` stand-in that keeps `len` and `shape` — legitimate scoring code allocates against them — while every conversion, index or attribute access raises a `RuntimeError` naming the rule. `load("test").aux` is similarly name-visible and value-hidden. Selection on the hidden split is thus not a policy the model is asked to respect but a capability it does not have.

The reported score. `SavedScoresEvaluator` (`loop.py`:104–190) requires the script to save the exact arrays behind its metrics, reloads them, checks shape, finiteness and length against the split, and recomputes `PRIMARY`, `GAUC` and `nDCG@5` with the pinned evaluator. The script’s own `METRICS` line is compared to that at a 10^{-6} tolerance — tight enough that a fabrication worth making is caught, loose enough that a script printing six significant figures is not. The tolerance was 10^{-9} and cost a legitimate iteration to a 5×10^{-7} display-rounding gap.

Validation reaching a fitted quantity. The critic’s fourth rung is the subtlest and was added last. An audit of 36 submitted scripts found two that fit weights on train but chose the *round count* by early stopping against `valid_sets=[<validation>]` — and then scored and reported those same validation rows. The critic now walks every call node for a `valid_sets`, `eval_set` or `validation_data` keyword whose value references a variable bound to `load("valid")`, and rejects it. What makes this rung necessary rather than redundant is that the effect is small: measured at +0.00041, roughly half a seed sigma, so it never trips the score-jump rung and nothing else in the chain was looking for it. The feedback returned to the proposer names the remedy — hold out the last days of *train* instead.

An externally checkable claim. `research/leakage_check.py` (28 lines) exists so that a reader does not have to take the preceding four layers on trust. It fits one `LightGBM` configuration

twice — train-only and train+validation — and prints both splits’ scores. If the agent were training on validation, its reported validation number would sit near the train+valid row rather than the train-only row. The submitted run’s 0.605309 is printed alongside for comparison. The design intent is stated in the module docstring: “*Run this yourself; trust the output, not us.*”

Leakage in the features themselves is handled a layer down. The post-click signals (`play_time_ms`, `is_click`, `is_like`, ...) are outcomes of the row being scored; `pipeline/data.py` exposes them only as `Split.aux`, never in `Split.X`, and the dataset’s supplied `video_features_statistic` file is not loaded at all because its counts aggregate over the full month and overlap the evaluation windows. `pipeline/history.py` recovers the same mechanism safely: entity histories aggregated over train rows only, leave-one-out on train so an example never receives its own target.

6 Search over solution scripts

Each scored script is a node in `tree.py` carrying its source, hypothesis, score, and counts of children, misses and crashes.

- **Miss.** A child failing to clear its parent by ε increments `misses`; two retire the node.
- **Crash.** A child that crashed or timed out counts half — the idea may be sound and the code wrong, and recovery already retries it — so four retire the node. Without this, a node that produced only crashes never approached retirement; one historical node spent five children and 78 minutes that way, one of them timing out at 4,373 s.
- **Selection.** `select("refine")` returns the highest-scoring live node; `select("sweep")` returns `min(live, key=(children, -score))` — the least-explored node, ties broken toward the higher score. Retirement plus this fallback is the backtracking a linear greedy walk cannot do.

`max_misses` defaults to 2 and must stay strictly below `patience`. When both were 3, checked against the same ε against the same reigning leader, a node was retired on exactly the iteration the run converged, and `select` never got to return the fallback — backtracking existed only for crash-heavy nodes.

Idea retirement. `recovery.py` keys an idea by the *method it names*, not the sentence, via an alias table (`lambdamart`, `ranknet`, `listwise` $\rightarrow$ `lambdarank`) and an overlap-coefficient match at 0.6. Overlap rather than Jaccard, because the model restates one idea at very different lengths and Jaccard reads a wordier rephrasing as a new idea. Two retries; six if the family has already scored well, on the reasoning that it is failing on code bugs rather than on being wrong. An idea that runs cleanly but loses to the incumbent by more than 2ε three times is retired too — crash-based blacklisting never fires for code that works.

7 Belief revision

After each scored iteration the agent rewrites a bounded set of claims, each with a status of `active`, `qualified` or `invalidated` and a list of supporting iteration numbers (`knowledge.py`). The point of the status field is that later evidence can *demote* an earlier conclusion rather than accumulating beside it; an append-only reflection log produced runs that recommended the same next experiment four times because nothing could overturn it. The rendered set — not a human-authored brief — is what guides the next proposal.

Two properties are enforced by construction: the set is capped at 14 claims so token use stays flat, and revision never raises. An LLM outage, a non-JSON reply or a malformed element leaves the previous belief set standing rather than ending an unattended run.

8 Portfolio mode

With `--slots  $n$` for $n \in \{2, \dots, 5\}$ the loop advances n lineages per turn. Figure 2 shows the turn structure.




Figure 2: One portfolio turn. Proposals are drawn one slot at a time (the LLM client mutates shared key-failover state), the scripts then run *concurrently*, and their results are processed back in slot order so a later slot can blend against an earlier one’s result. The organizers’ stopping rule is singular — there is no per-agent counter in it — so n lineages cannot each carry one. What they *can* do is share a turn: the turn’s score is the best of its slots, and that single curve is what the rule is measured against. Stale slots are archived and refilled between turns.

The justification for the whole portfolio is falsifiable and is measured every turn. `pairwise_rank_correlation` computes within-user rank correlation between the slots’ pre-blend validation predictions; ranks are taken inside each user because the metric only ever compares rows belonging to the same user. If that number sits above 0.95, the slots are ranking validation identically and the extra cost bought nothing. Measuring it *pre-blend* matters: one run read 1.000 because two slots had blended to $\alpha = 0$ and were literally holding the same incumbent array.

9 Results: run r96

`runs/r96` is the submitted run. It was produced under the `train-only-v3` data contract, reached the organizers’ convergence rule in 14 of the 50 permitted iterations, and required no human intervention of any kind. Every figure below is re-derived from that run’s own records — `ledger.jsonl`, `llm_calls.jsonl`, `candidates.jsonl`, `portfolio.jsonl`, `archive.jsonl` and `harness_ensembles.jsonl` — all of which ship with the repository.

9.1 Headline

The run reports no hidden-test score of its own, by design. `test_scored` is `false` in `run_meta.json`: the harness writes the submission *without* evaluating it, because evaluating it would mean reading test labels, which §2.9.3 forbids in any form. The hidden-test figure quoted

	validation primary	delta vs. official baseline
Official baseline (organizer-published)	0.6016	—
Agent’s own baseline reproduction (iter. 1)	0.602009	+0.000409
Validation-best checkpoint (iter. 11)	0.605759	+0.004159

Table 3: The scored quantity is the validation-best checkpoint at the point the convergence rule fired. The agent’s own reproduction lands 0.000409 from the organizers’ published 0.6016 — comfortably inside the 0.002 tolerance — which is what satisfies Task Requirement 1 before any improvement is attempted.

in README.md — GAUC 0.6668 / nDCG@5 0.5315 / PRIMARY 0.59912, a delta of +0.00452 — was produced afterwards, outside the run, by scoring the submitted file once. Keeping that step outside the loop is what makes the compliance position in Section 11 clean.

9.2 Trajectory

Fourteen scripts across four improve turns with three parallel slots: zero failures, zero integrity rejections, zero manual interventions.

iter	turn	slot	phase	status	primary	sec
0	0	—	eda	ok	—	28
1	0	—	baseline	ok	0.602009	144
2	1	0	improve	ok	0.604169	783
3	1	1	improve	kept	0.604654	826
4	1	2	improve	kept	0.605036	1042
5	2	0	improve	kept	0.605584	161
6	2	1	improve	kept	0.605584	74
7	2	2	improve	kept	0.605584	23
8	3	0	improve	kept	0.605607	309
9	3	1	improve	kept	0.605607	353
10	3	2	improve	kept	0.605633	25
11	4	0	improve	kept	0.605759	340
12	4	1	improve	kept	0.605759	374
13	4	2	improve	kept	0.605759	35

Table 4: Every scored iteration. **kept** means the gain did not clear ε ; the score is still ranked and still submittable, and the submitted checkpoint came from such an iteration. The PRIMARY column is the retained score after the controller’s blend (Section 9.4), which is why the curve is monotone.

9.3 Convergence

The run stopped on the organizers’ default rule — $\varepsilon = 0.002$, $N = 3$, no declared minimum-iteration floor — and the stop reproduces exactly from the ledger. Rebuilding the best-validation curve and applying `agent.loop.converged`, the rule itself rather than a second copy of it:

turn	scripts	turn best	running best	converged?
1	3	0.605036	0.605036	no
2	3	0.605584	0.605584	no
3	3	0.605633	0.605633	no
4	3	0.605759	0.605759	yes

At turn 4 the window gain is $0.605759 - 0.605036 = 0.000723 \leq \varepsilon$, and the run stops. This follows FAQ 2.9.1’s cumulative wording — the best over the last N scored steps against the best from before that window — rather than requiring each individual step to clear ε .

Because a portfolio turn launches three scripts, the run reports both readings of the iteration budget: `scripts` is 14 and `turns` is 4, so it can be checked against either. Under the strictest per-script reading the rule would have fired at iteration 5 with a validation best of 0.605584; continuing to turn 4 added 0.000175, about a fifth of the baseline’s own 5-seed σ of 0.0008. Fourteen scripts sits inside the 50-iteration cap under either reading, and both figures are in the run log.

9.4 Where the gain came from

The most productive mechanism in the run is the controller’s α -blend against the trusted incumbent (Section 5) — chosen on validation, applied unchanged to test. Four of the twelve improve iterations produced a candidate that scored *at or below* the standing incumbent and still advanced the leader, because a mixture of the two beat both:

iter	candidate	incumbent	selected α	retained	effect
2	0.604059	0.602009	0.75	0.604169	blend beats both
3	0.604516	0.604169	0.75	0.604654	blend beats both
4	0.604332	0.604654	0.25	0.605036	candidate <i>loses</i> , blend still gains
5	0.605584	0.605036	1.00	0.605584	candidate taken outright
6	0.604796	0.605584	0.00	0.605584	candidate declined, incumbent held
7	0.605137	0.605584	0.00	0.605584	candidate declined, incumbent held
8	0.605606	0.605584	0.75	0.605607	marginal
9	0.605577	0.605607	0.05	0.605607	marginal
10	0.605605	0.605607	0.50	0.605633	candidate <i>loses</i> , blend still gains
11	0.605720	0.605633	0.50	0.605759	submitted checkpoint
12	0.605677	0.605759	0.00	0.605759	candidate declined
13	0.605658	0.605759	0.00	0.605759	candidate declined

Table 5: Every harness blend decision, from `harness_ensembles.jsonl`. $\alpha = 0$ retains the incumbent, $\alpha = 1$ takes the candidate outright. The grid is $\{0.05, 0.10, 0.15, 0.20, 0.25, 0.50, 0.75\}$ plus both endpoints, and a weight enters the full-window comparison only if it is non-negative on *both* user-grouped validation folds — so a weight that wins one half of validation and loses the other is never chosen.

Two properties make this safe as well as effective. It is a controller-side decision over predictions the trusted evaluator has already verified, so the agent is never grading itself; and the fold-A/fold-B admission rule means a blend weight has to survive a held-out half of validation before it can be selected on the whole of it.

9.5 Portfolio behaviour

Three slots advanced per turn. The measured within-user rank correlation between the slots’ *pre-blend* predictions — the acceptance test for the whole portfolio design — shows the lineages separating as the run progresses:

turn	mean	max	pairs (0–1, 0–2, 1–2)
1	0.9017	0.9188	0.878, 0.908, 0.919
2	0.6509	0.7327	0.733, 0.585, 0.635
3	0.7397	0.9111	0.911, 0.654, 0.654
4	0.5676	0.7545	0.754, 0.437, 0.512

No turn tripped the 0.95 redundancy alarm, and mean agreement fell from 0.90 to 0.57 across the run: the three lineages were genuinely exploring different parts of the space rather than rediscovering one another’s answers. All three slots were archived at turn 3 with their predictions retained, and the sibling-disclosure mechanism (Section 8) is the most likely reason the separation held.

9.6 Resource usage

Iterations / scripts	14 of the 50 cap	
Turns (3 slots each)	4	
Agent wall-clock	53.0 min of the 6 h ceiling	
Summed script time	75.3 min (1.42× wall-clock; slots run concurrently)	
LLM calls	19	
Tokens in / out	167,466 / 93,501	
Tokens total	260,967	
GPU-hours	0 — CPU only	
Manual interventions	0	
Failures / integrity rejections	0 / 0	
Candidates compared	316 models inside those 14 scripts	

Table 6: The resource figures required by Deliverable 4. The 19 LLM calls decompose, from `llm_calls.jsonl`, as 1 EDA proposal + 1 baseline proposal + 1 belief revision + 13 improve proposals + 3 cross-slot consultant syntheses. Thirteen proposals for twelve improve scripts: one reply did not satisfy the response contract and was re-requested, which is the proposer’s built-in retry rather than an intervention. The model is never called inside a training loop.

10 Findings

10.1 One iteration is one script, not one model

This is the design decision the result rests on. The convergence rule charges per *iteration*, so a script that builds and compares a dozen models costs exactly what a script that builds one costs. **r96 evaluated 316 distinct models across 14 charged iterations** — a 22× multiplier on the budget the rule actually meters.

The agent discovered this and pressed on it. Candidates compared per turn rose monotonically as the run went on:

turn	candidates	best of turn
1	58	0.604516
2	66	0.605584
3	86	0.605606
4	106	0.605720

Nothing in the prompt sets a target number. The agent widened its internal search itself, having seen from its own history that breadth inside a script is free while breadth across iterations is what ends the run.

10.2 The agent found the binding constraint before it modelled anything

The EDA script (iteration 0) is written by the agent, and what it prints is the only dataset knowledge the run ever has. It measured the distribution shift across the split boundary without being told to look for one:

	train	validation	
positive rate	0.3366	0.3133	falling
median rows per user	31	4	8× sparser
users with no positive label	5.1%	30.3%	6× more

It also printed the per-day label rate across the validation week, showing a monotone decline ($0.319 \rightarrow 0.290$), and measured that 7.7% of adjacent same-user impressions share a timestamp — a fact about how the feed is logged that bears directly on whether sequence features are constructible.

Every winning hypothesis that follows targets that drift. Iteration 11, the submitted checkpoint, reads:

*At the objective and sequence-feature stages, compare recency-weighted LambdaMART, setwise boosted **rank_xendcg**, and a non-parametric chronological transition model; ranking losses directly emphasize within-user ordering while transition rates exploit short-term feed context that static identity models miss.*

The chain from measurement to target to method is the agent’s own: it measured the drift, named it as the stage to attack, and selected the methods.

10.3 Breadth across method families, not tuning

The 316 candidates span the recommender literature rather than one family’s hyperparameters. The strongest candidate in each turn:

turn	best candidate	primary
1	dcn	0.604516
2	din_blend_0.40	0.605584
3	fibinet_blend_0.15	0.605606
4	recency_rank_xendcg_incbblend_0.20	0.605720

Across the run the agent built and scored DCN, DeepFM, xDeepFM, FiBiNET, PNN, FFM, DIN, MMoE, truncated SVD, LambdaMART, **rank_xendcg**, graph-diffusion and non-parametric transition models — and, in the later turns, cross-family compositions such as **pnn_fibinet_equal_incbblend_0.35**. That progression, from single architectures in turn 1 to composed ensembles of decorrelated families by turn 4, is the search behaviour the design is meant to produce.

10.4 Robustness: fourteen scripts, zero failures

Across 14 generated scripts totalling roughly 200 KB of model-written code — building 316 models on 1.14M training rows under a per-script timeout — the run recorded **zero** crashes, zero timeouts, zero integrity rejections and zero manual interventions.

Two harness decisions are visible behind that number. Generated scripts run with LightGBM preloaded ahead of torch, which removes an OpenMP ordering clash that previously cost whole iterations to segfaults and hangs; and the trusted evaluator’s 10^{-6} agreement tolerance is loose enough that an honest script printing six significant figures is not rejected, while still being $2000\times$ tighter than the ε that decides convergence.

10.5 Efficiency

The run reached its converged result on 28% **of the iteration budget** (14 of 50) and 15% **of the wall-clock ceiling** (53 min of 6 h), on CPU alone, for 260,967 tokens. Summed script time exceeds wall-clock by $1.42\times$, which is the portfolio doing its job: three lineages advancing concurrently inside the same elapsed hour.

11 Compliance with the problem statement

The organizers state that constraints are checked by review of both the prediction file and the code and run logs (§2.9.3). The system is built so that review is straightforward, and in several places the constraint is enforced by construction rather than by instruction.

Constraint	How it is satisfied	Evidence
§2.2/1 Reproduce the official baseline	Iteration 1 stands up an end-to-end FM pipeline reaching 0.602009 against the published 0.6016. The controller accepts only within ± 0.002 and otherwise feeds the error back and retries.	<code>ledger.jsonl</code> iter. 1
§2.9.2 Train split only	No model whose predictions are saved is fitted on validation. AST analysis of all 14 scripts finds no train+validation concatenation, no <code>valid_sets/eval_set</code> /early stopping against validation, and no feature statistic computed over validation rows.	Section 12; <code>critic.py</code> <code>eval_set</code> rung
§2.9.2 No <code>log_random</code>	The loader never opens it — splits are built only from the two standard logs. Zero references in any generated script.	<code>data.py</code> <code>_TRAIN_LOG</code> , <code>_VALID_TEST_LOG</code>
§2.9.2 No 1k/27k as auxiliary	<code>KUAI RAND_VARIANT</code> binds one release per cache; <code>r96</code> ran on <code>KuaiRand-Pure</code> alone.	<code>dataset</code> in <code>run_meta.json</code>
§2.9.3 No test labels, in any form	Enforced as a <i>capability</i> , not a rule: under <code>AGENT_HIDE_TEST_LABELS=1</code> the loader returns a stand-in whose every value access raises, so a script cannot read a test label even by accident. AST analysis finds no <code>.y</code> or <code>.aux</code> access on a test-bound variable anywhere. The harness writes the submission without scoring it.	<code>data.py</code> <code>HiddenTestLabels</code> , <code>HiddenTestOutcomes</code> ; <code>test_scored: false</code>
§2.9.1 Convergence rule	The organizers’ defaults are declared and recorded: $\epsilon = 0.002$, $N = 3$, no minimum-iteration floor. The stop reproduces from the ledger.	<code>epsilon</code> , <code>patience</code> , <code>min_iters</code> ; Section 9
Limits: 50 iterations	14 scripts executed. Both readings recorded — <code>scripts</code> (14) and <code>turns</code> (4) — so the run is checkable against either.	<code>iterations</code> , <code>scripts</code> , <code>turns</code>
Limits: 6 h wall-clock	53.0 min. Measured as true elapsed controller time, not summed script time, and enforced as a stop condition inside the loop.	<code>wall_clock_s</code>
No external training data	Only the <code>KuaiRand-Pure</code> cache is reachable from the loader; no pretrained weights are used.	<code>data.py</code> ; generated scripts
Submission schema	<code>row_id, user_id, video_id, score</code> ; 170,588 data rows matching the test split, all scores finite.	<code>submission_best.csv</code>
Deliverable 3: run logs	Per-iteration hypothesis, complete script, metrics, status and recovery events, plus every LLM prompt and response.	<code>ledger.jsonl</code> , <code>scripts/</code> , <code>llm_calls.jsonl</code>

Table 7: Constraint-by-constraint position for `r96`.

11.1 Three clarifications, and what following them cost

Three FAQ clarifications landed against decisions this project had originally made the other way. Each was reversed, and the cost of reversing it is recorded rather than absorbed quietly.

- **§2.9.2 — training data is the train split only.** Earlier runs refit the final model on train+validation before predicting test. Measured on the same harness, that reading was

worth 0.00229 of hidden-test delta. It is removed, and `critic.py` now rejects any iteration that fits on validation, including through early stopping or feature statistics. The reported result is the smaller, compliant one.

- **§2.9.3 — test labels may not be used in any way.** Test labels are hidden by default in the loader, and the harness writes the submission without scoring it. The hidden-test figure was produced once, after the run, outside the loop.
- **§2.9.1 — a team may declare its own stopping rule.** We declare the organizers' defaults, $\varepsilon = 0.002$ and $N = 3$ with no minimum-iteration floor, and record all three in `run_meta.json`. A floor that kept searching past the stop point was tested and dropped: it bought +0.00047 of validation and nothing on test.

One record-keeping note belongs with the third point, since a reviewer comparing `run_meta.json` against `README.md` will meet it. The run was launched with that exploratory floor still enabled, continued past the declared stop, and was halted; the retained record is truncated to the point where the declared no-floor rule fires, and `run_meta.json` was rebuilt from the surviving logs (its `reconstructed_note` says so). The truncation is in the conservative direction — the declared rule stops *earlier* and yields a *lower* validation best than the run actually reached — so it cannot flatter the result. Relaunching with `--min-iters 0` would make the declaration and the launch the same object, and would not change the reported numbers.

12 Verification and reproducibility

The run's claims were checked against its own records, independently of the code that produced them. Everything below was confirmed.

- **The metadata is faithful to the raw records.** Because `run_meta.json` was reconstructed, all sixteen of its derived fields were recomputed from `ledger.jsonl`, `llm_calls.jsonl` and `candidates.jsonl` — iterations, scripts, turns, slots, failures, integrity rejections, summed script seconds, tokens in/out/total, candidates evaluated, EDA completion, the baseline score, and the submitted iteration's id and validation primary. Every field matched, with no discrepancy.
- **The convergence stop reproduces.** Rebuilding the best-validation curve from the ledger and applying `agent.loop.converged` — imported from the controller, not reimplemented — fires at turn 4 with a window gain of 0.000723, exactly as recorded.
- **No script touches a test label or an evaluation-split outcome.** An AST pass over all 14 scripts resolved every variable bound to `load("train"/"valid"/"test")` and searched for `.y` or `.aux` access on test-bound variables, `.aux` on validation-bound variables, and validation-bound variables reaching a fitting call. **Zero findings.** The repository's own `critic.py` independently flags zero of the 14.
- **No forbidden source is referenced.** No occurrence of `log_random`, `video_features_statistic`, or any train+validation concatenation. The only `.aux` uses are two references to `train.aux` in `iter_4.py` as auxiliary supervision targets — explicitly legitimate, and a use the critic is written to permit rather than flag.
- **The submission file is well-formed.** 170,588 data rows matching the test split size, correct header, all scores finite.
- **The test suite passes.** All 25 modules.

12.1 Provenance of the scores

The validation figures throughout this report are read from the ledger rather than recomputed from the dataset. That chain is strong by construction: a score reaches the ledger only after `SavedScoresEvaluator` has reloaded the script’s saved predictions and recomputed the metric with the pinned evaluator, and only if the script’s self-reported figure agrees to 10^{-6} . The number in the ledger is the trusted evaluator’s, never the script’s own claim.

For an external check that needs no trust in this harness at all, the repository ships `research/leakage_check.py`: it fits one configuration twice, train-only and train+validation, and prints both splits, so the submitted validation figure can be placed against the two. `research/verify_claims.py` re-derives the reported tables from the run records and exits non-zero on any disagreement. Both require the full KuaiRand-Pure cache.

13 Reproducing a run

```
# data (194 MB, public)
mkdir -p data/raw && cd data/raw
curl -L0 https://zenodo.org/records/10439422/files/KuaiRand-Pure.tar.gz
tar -xzf KuaiRand-Pure.tar.gz && cd ../../
python -c "import pipeline.data as d; \
          d.build_cache('data/raw/KuaiRand-Pure/data', 'data/cache')"

export OPENAI_API_KEY=...          # and LLM_PROVIDER=openai, or ANTHROPIC_API_KEY

python -m pipeline.baseline_fm      # the official baseline, for reference
python run_agent.py --run-dir runs/rN --iters 50 --timeout 900 --patience 3
python report_run.py runs/rN > RUN_REPORT.md
python -m pipeline.submit --check --in runs/rN/submission.csv --split test

python -m research.leakage_check    # train-only vs train+valid, side by side
python -m research.verify_claims    # re-derive the ablation tables
```

Cross-run memory is off by default; pass `--memory` to opt in. `--slots n` selects portfolio mode. `--force-mode` pins the search mode for a diagnostic run and is not for scoring.

Two offline modes exist. `--dry-run` exercises the entire loop with a canned LLM and no network. `--replay runs/rN/llm_calls.jsonl` re-runs a previous run against its recorded responses: deterministic, no tokens, and it reproduces the original ledger to the last decimal. Replay tests the loop, the parsers, the ledger and the reporting — *not* prompts, since a changed prompt still receives the response recorded for the old one. `--replay-strict` fails the moment a prompt diverges, so that limitation cannot pass unnoticed.